



PLATFORM SETUP GUIDES / [scheduler-setup-guide](#)

# Contents

- PLATFORM SETUP GUIDES / Scheduler setup guide
  - Introduction
  - Infrastructure prerequisites
  - Dependencies
  - Dependencies
    - MongoDB helm example
  - Configuration
    - Configuring MongoDB
    - Configuring Kafka
    - Configuring logging

## PLATFORM SETUP GUIDES / Scheduler setup guide

### Introduction

This guide will walk you through the process of setting up the Scheduler service using a Docker image.

### Infrastructure prerequisites

- **MongoDB** - version 4.4 or higher for storing taxonomies and contents
- **Kafka** - version 2.8 or higher

### Dependencies

- **MongoDB** database
- ability to connect to a Kafka instance used by the engine

The service comes with most of the needed configuration properties filled in, but there are a few that need to be set up using some custom environment variables.

## Dependencies

### MongoDB helm example

Basic MongoDB configuration - helm values.yaml

```
scheduler-mdb:
  existingSecret: {{secretName}}
  mongodbDatabase: {{SchedulerDatabaseName}}
  mongodbUsername: {{SchedulerDatabaseUser}}
  persistence:
    enabled: true
    mountPath: /bitnami/mongodb
    size: 4Gi
  replicaSet:
    enabled: true
    name: rs0
    pdb:
      enabled: true
      minAvailable:
        arbiter: 1
        secondary: 1
  replicas:
    arbiter: 1
    secondary: 1
```

```
useHostnames: true
serviceAccount:
  create: false
usePassword: true
```

### DANGER

This service needs to connect to a Mongo database that has replicas, in order to work correctly.

## Configuration

### Configuring MongoDB

The MongoDB database is used to persist scheduled messages until they are sent back. The following configurations need to be set using environment variables:

- `SPRING_DATA_MONGODB_URI` - the URI for the MongoDB database

### Configuring Kafka

The following Kafka related configurations can be set by using environment variables:

- `SPRING_KAFKA_BOOTSTRAP_SERVERS` - address of the Kafka server
- `SPRING_KAFKA_CONSUMER_GROUP_ID` - group of consumers
- `KAFKA_CONSUMER_THREADS` - the number of Kafka consumer threads

- `KAFKA_AUTH_EXCEPTION_RETRY_INTERVAL` - the interval between retries after `AuthorizationException` is thrown by `KafkaConsumer`

Each action available in the service corresponds to a Kafka event. A separate Kafka topic must be configured for each use-case.

### CAUTION

Make sure the topics configured for this service don't follow the engine pattern.

## Configuring logging

The following environment variables could be set in order to control log levels:

- `LOGGING_LEVEL_ROOT` - root spring boot microservice logs
- `LOGGING_LEVEL_APP` - app level logs

**Was this page helpful?**